

---

# Algorithms for Massive, Expensive, and Otherwise Inconvenient Graphs

---

DAVID TENCH

UNIVERSITY OF MASSACHUSETTS AMHERST




---

Ask questions if you're confused!

---

Ask questions if you're confused!

Much of this presentation requires  
basic CS knowledge only.

A solid green horizontal bar at the bottom of the slide.

---

Ask questions if you're confused!

Much of this presentation requires  
basic CS knowledge only.

I'll warn you before the tricky parts.

A solid green horizontal bar at the bottom of the slide.

---

# Algorithms for Massive, Expensive, and Otherwise Inconvenient Graphs

---

DAVID TENCH

UNIVERSITY OF MASSACHUSETTS AMHERST




# Convenient Inputs

---

# Convenient Inputs

---

- What is the 4<sup>th</sup> element of this list?





8	6	1	9	3	4	5	2
---	---	---	---	---	---	---	---

# Convenient Inputs

---

- What is the 4<sup>th</sup> element of this list?
- Is there an edge between some pair of nodes in this graph?







# Convenient Inputs

- What is the 4<sup>th</sup> element of this list?
- Is there an edge between some pair of nodes in this graph?

An algorithm can access **any part** of its input at **any time** at **unit cost**.





# Convenient Inputs

- What is the 4<sup>th</sup> element of this list?
- Is there an edge between some pair of nodes in this graph?

An algorithm can access **any part** of its input at **any time** at **unit cost**.

This is the *random access property*.



# When Random Access Fails

---

# When Random Access Fails

---

When inputs are too large  
to fit in memory

# When Random Access Fails

---

When inputs are too large  
to fit in memory



# When Random Access Fails

---

When inputs are too large  
to fit in memory

When parts of inputs are  
costly to discover




# When Random Access Fails

---

When inputs are too large  
to fit in memory



When parts of inputs are  
costly to discover



# Overview of this talk

---

Streaming graph algorithms:  
Coping with graphs too large to  
fit in memory






# Overview of this talk

# Streaming graph algorithms: Coping with graphs too large to fit in memory



## Collaborations with practitioners: Graph algorithms subject to expensive edge queries




# Overview of this talk

Streaming graph algorithms:  
Coping with graphs too large to  
fit in memory



Collaborations with practitioners:  
Graph algorithms subject to  
expensive edge queries



(Also future work)

---

# Streaming

---

COMPUTING WITH INCREDIBLY LARGE INPUTS



# When Graphs Are Too Large

## Can't store graph in memory



# When Graphs Are Too Large

## Can't store graph in memory

## Receive *stream* of edges





# When Graphs Are Too Large

## Can't store graph in memory

## Receive *stream* of edges

- Vertex Connectivity (PODS 2015)
- Densest Subgraph (MFCS 2015)
- Unique Cover (current work)




# Warm-up: Missing Number

---

I read you an unordered list of all the integers from 1 to 5 million – except I leave one of them out. After, I ask you which is missing.

You only have a single piece of paper and a pencil. How do you find the missing number?




# Warm-up: Missing Number

---

I read you an unordered list of all the integers from 1 to 5 million – except I leave one of them out. After, I ask you which is missing.

You only have a single piece of paper and a pencil. How do you find the missing number?

Answer: keep a running sum of all the numbers, then subtract from  $1 + 2 + \dots + 5 \text{ million}$ .





# Graph Streaming

---



# Graph Streaming

Facebook has almost 2 billion users



# Graph Streaming

# Facebook has almost 2 billion users

Nodes = users, edges = friend relationships



# Graph Streaming



Facebook has almost 2 billion users

Nodes = users, edges = friend relationships

This graph could have almost  $(2 \text{ billion})^2 = 4 \text{ quintillion}$  edges



# Graph Streaming



# Facebook has almost 2 billion users

Nodes = users, edges = friend relationships

This graph could have almost  $(2 \text{ billion})^2 = 4 \text{ quintillion edges}$

For modern computers, storing  
2 billion objects is maybe  
reasonable but 4 quintillion is not

# Graph Streaming



It's possible to store some of the edges, but not all

# Graph Streaming



It's possible to store some of the edges, but not all

We have roughly  $n$  space,  
where  $n = \#$  of nodes

# Graph Streaming



It's possible to store some of the edges, but not all

We have roughly  $n$  space,  
where  $n = \#$  of nodes

## Graph edges are given as a *stream*



# Defining a graph via a stream

---



# Defining a graph via a stream

---

Add edge (1,4)



# Defining a graph via a stream

---



# Defining a graph via a stream


---

Add edge (3,4)



# Defining a graph via a stream


---



# Defining a graph via a stream


---

Add edge (1,3)



# Defining a graph via a stream


---



# Defining a graph via a stream

---


Add edge (2,5)





# Defining a graph via a stream

---




# Defining a graph via a stream

---

Stream:

Add edge (1,4)  
Add edge (3,4)  
Add edge (1,3)  
Add edge (2,5)

Resulting Graph:




# Defining a graph via a stream

---

Stream:

Add edge (1,4)  
Add edge (3,4)  
Add edge (1,3)  
Add edge (2,5)


Resulting Graph:




Edge deletions are also possible,  
but we're ignoring them today.

# Problem: Vertex Connectivity

---




What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?




# Problem: Vertex Connectivity

---






What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?



# Problem: Vertex Connectivity




---



What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?

# Problem: Vertex Connectivity




---



What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?

# Problem: Vertex Connectivity

---






What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?

How many computers in a network can fail before the remaining network is disconnected?



# Problem: Vertex Connectivity

---




What is the minimum number of nodes we can remove to disconnect a graph  $G$ ?


Easily solvable using max flow algorithm, but this requires random access to the graph.

# Problem: Vertex Connectivity

---



G




H


We show how to create a *certificate* graph H that matches G's vertex connectivity up to constant  $k$ , but has only roughly  $kn$  edges.

# Problem: Vertex Connectivity

---



G




H

We show how to create a *certificate* graph  $H$  that matches  $G$ 's vertex connectivity up to constant  $k$ , but has only  $kn$  edges.


We also show how to  $(1+\epsilon)$ -approximate vertex connectivity while only storing  $\epsilon^{-1} kn$  edges.

# Problem: Vertex Connectivity

---



G



H


We show how to create a *certificate* graph  $H$  that matches  $G$ 's vertex connectivity up to constant  $k$ , but has only  $kn$  edges.

We also show how to  $(1+\epsilon)$ -approximate vertex connectivity while only storing  $\epsilon^{-1} kn$  edges.


Finally, we show how to construct hypergraph sparsifiers in roughly linear space.

# Problem: Vertex Connectivity

---



G



H

We show how to create a *certificate* graph  $H$  that matches  $G$ 's vertex connectivity up to constant  $k$ , but has only  $kn$  edges.


We also show how to  $(1+\epsilon)$ -approximate vertex connectivity while only storing  $\epsilon^{-1} kn$  edges.

Finally, we show how to construct hypergraph sparsifiers in roughly linear space.


# Warm-up: Is the Graph Connected?

---

Original  
graph




Summary



# Warm-up: Is the Graph Connected?


---

Original  
graph




As each edge arrives in the stream, we keep it only if its endpoints were not already connected.

Summary




# Warm-up: Is the Graph Connected?

---



As each edge arrives in the stream, we keep it only if its endpoints were not already connected.






# Warm-up: Is the Graph Connected?


---

Original  
graph



As each edge arrives in the stream, we keep it only if its endpoints were not already connected.


Summary



# Warm-up: Is the Graph Connected?


---

Original  
graph



As each edge arrives in the stream, we keep it only if its endpoints were not already connected.



Summary



# Warm-up: Is the Graph Connected?



---

Original  
graph



As each edge arrives in the stream, we keep it only if its endpoints were not already connected.



Summary



# Warm-up: Is the Graph Connected?



---

Original  
graph



As each edge arrives in the stream, we keep it only if its endpoints were not already connected.

Spanning  
Forest





This is a **spanning forest** and has at most  $n-1$  edges.



# Warm-up: Is the Graph Connected?

---

Original  
graph



Spanning  
Forest




As each edge arrives in the stream, we keep it only if its endpoints were not already connected.

This is a **spanning forest** and has at most  $n-1$  edges.

If the spanning forest is connected, we know the original graph was as well.


# Back to Vertex Connectivity

---



# Back to Vertex Connectivity


---



First we delete each node in  $G$  with probability  $1-1/k$ .

# Back to Vertex Connectivity

---




First we delete each node in  $G$  with probability  $1-1/k$ .



# Back to Vertex Connectivity

---




First we delete each node in  $G$  with probability  $1-1/k$ .

During the stream we maintain a spanning forest  $T$  on the remaining nodes.

# Back to Vertex Connectivity

---




First we delete each node in  $G$  with probability  $1-1/k$ .

During the stream we maintain a spanning forest  $T$  on the remaining nodes.

# Back to Vertex Connectivity

---




First we delete each node in  $G$  with probability  $1-1/k$ .

During the stream we maintain a spanning forest  $T$  on the remaining nodes.

# Back to Vertex Connectivity


---



**G**


First we delete each node in  $G$  with probability  $1-1/k$ .

During the stream we maintain a spanning forest  $T$  on the remaining nodes.






**T**

# Back to Vertex Connectivity



First we delete each node in  $G$  with probability  $1-1/k$ .


During the stream we maintain a spanning forest  $T$  on the remaining nodes.




Repeat roughly  $k^2$  times in parallel\* to make  $T_1, T_2, \dots$

# Back to Vertex Connectivity

---




After the stream, merge all  $T_i$  to form  $H$ , our vertex connectivity certificate.




# Back to Vertex Connectivity

---





After the stream, merge all  $T_i$  to form  $H$ , our vertex connectivity certificate.



# Back to Vertex Connectivity

---





After the stream, merge all  $T_i$  to form  $H$ , our vertex connectivity certificate.

Each spanning forest has  $n/k$  edges in expectation and there are  $k^2$  of them so  $H$  has  $kn$  edges in expectation.



# Back to Vertex Connectivity

---



After the stream, merge all  $T_i$  to form  $H$ , our vertex connectivity certificate.

Theorem: For some arbitrary set  $S$  of at most  $k$  nodes,  $G \setminus S$  is disconnected iff  $H \setminus S$  is disconnected.

# Proof Sketch of Theorem

---

Theorem: For some arbitrary set  $S$  of at most  $k$  nodes,  $G \setminus S$  is disconnected iff  $H \setminus S$  is disconnected.

# Proof Sketch of Theorem

---

To prove: When  $G \setminus S$  is disconnected,  $H \setminus S$  must be disconnected

To prove: When  $G \setminus S$  is connected,  $H \setminus S$  must be connected

# Proof Sketch of Theorem

---


To prove: When  $G \setminus S$  is disconnected,  $H \setminus S$  must be disconnected

To prove: When  $G \setminus S$  is connected,  $H \setminus S$  must be connected

# Proof Sketch of Theorem

---

To prove: When  $G \setminus S$  is disconnected,  $H \setminus S$  must be disconnected




# Proof Sketch of Theorem

---

To prove: When  $G \setminus S$  is disconnected,  $H \setminus S$  must be disconnected

$H$  is a subgraph of  $G$ , so if  $G \setminus S$  lacks a path between nodes  $u$  and  $v$ , so does  $H \setminus S$ .




# Proof Sketch of Theorem

---

To prove: When  $G \setminus S$  is disconnected,  $H \setminus S$  must be disconnected

$H$  is a subgraph of  $G$ , so if  $G \setminus S$  lacks a path between nodes  $u$  and  $v$ , so does  $H \setminus S$ .

With high probability,  $H$  has all of  $G$ 's nodes. //



# Proof Sketch of Theorem

---

To prove: When  $G \setminus S$  is connected,  
 $H \setminus S$  must be connected




# Proof Sketch of Theorem

---

To prove: ~~When  $G \setminus S$  is connected,~~  
 ~~$H \setminus S$  must be connected~~

To prove: If some edge  $(u,v)$  exists  
in  $G \setminus S$ , then there is a path  
between  $u$  and  $v$  in  $H \setminus S$ .


# Proof Sketch of Theorem



To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

# Proof Sketch of Theorem

---




To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

If there's some  $T_i$  that contains both  $u$  and  $v$ ...

# Proof Sketch of Theorem

---



**G**

To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .


If there's some  $T_i$  that contains both  $u$  and  $v$ ...

Either edge  $(u,v)$  is in  $T_i$  or

**$T_i$**

# Proof Sketch of Theorem

---



**G**

To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

If there's some  $T_i$  that contains both  $u$  and  $v$ ...


Either edge  $(u,v)$  is in  $T_i$  or

Some other path between  $u$  and  $v$  is in  $T_i$

**$T_i$**

# Proof Sketch of Theorem

---



**$G$**

To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .


If there's some  $T_i$  that contains both  $u$  and  $v$ ...

**$T_i$**

The only potential problem is that this alternate path may contain a node that's in  $S$  (red node).

# Proof Sketch of Theorem

---



**G**

To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .


If there's some  $T_i$  that contains both  $u$  and  $v$ ...

**T<sub>i</sub>**

The only potential problem is that this alternate path may contain a node that's in  $S$  (red node).

So we need  $T_i$  to not contain any nodes in  $S$ .

# Proof Sketch of Theorem




To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

We need  $T_i$  to contain  $u$  and  $v$ , and no red  $S$  nodes.

$$P(u \text{ and } v \text{ connected in } T_i \setminus S) = \frac{1}{k^2} \left(1 - \frac{1}{k}\right)^k$$



# Proof Sketch of Theorem




To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

We need  $T_i$  to contain  $u$  and  $v$ , and no red  $S$  nodes.

$$P(u \text{ and } v \text{ disconnected in } T_i \setminus S) = 1 - \frac{1}{k^2} \left(1 - \frac{1}{k}\right)^k$$

# Proof Sketch of Theorem



**G**

To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .

We need  $T_i$  to contain  $u$  and  $v$ , and no red  $S$  nodes.


$$P(u \text{ and } v \text{ disconnected in } T_i \setminus S) = 1 - \frac{1}{k^2} \left(1 - \frac{1}{k}\right)^k$$

**T<sub>i</sub>**

$$P(u \text{ and } v \text{ disconnected in } H \setminus S) = \left(1 - \frac{1}{k^2} \left(1 - \frac{1}{k}\right)^k\right)^{O(k^2 \log(n))} \leq \frac{1}{n^4}$$

# Proof Sketch of Theorem

---



To prove: If edge  $(u,v)$  exists in  $G \setminus S$ ,  $u$  and  $v$  are connected in  $H \setminus S$ .


We need  $T_i$  to contain  $u$  and  $v$ , and no red  $S$  nodes.

The probability that no  $T_i$  meets this requirement is at most  $1/n^4$ .


So  $u$  and  $v$  are connected in  $H \setminus S$  with high probability. //

# We did it!

---



G




H


We showed how to create a *certificate* graph H that matches G's vertex connectivity up to constant  $k$ , but has only roughly  $kn$  edges.

# We did it!

---



G



H

We showed how to create a *certificate* graph H that matches G's vertex connectivity up to constant  $k$ , but has only roughly  $kn$  edges.

You can process a massive stream of the edges in G to create H, which is much smaller. Then you can run a traditional vertex connectivity algorithm on H to get your answer.

---

# Query-Based Algorithms

---

WHEN DISCOVERING GRAPH EDGES IS COSTLY



# When Graph Edges Are Costly

---

Can check existence of any edge at any time, but pay a significant cost

Want to minimize # of queries


# When Graph Edges Are Costly

---

Can check existence of any edge at any time, but pay a significant cost

Want to minimize # of queries

Mesh memory manager (PLDI 2019)






# When Graph Edges Are Costly


Can check existence of any edge at any time, but pay a significant cost

Want to minimize # of queries

Mesh memory manager (PLDI 2019)



PathCache network path predictor  
(to be submitted SIGCOMM 2020)




# When Graph Edges Are Costly


Can check existence of any edge at any time, but pay a significant cost

Want to minimize # of queries

Mesh memory manager (PLDI 2019)



PathCache network path predictor  
(to be submitted SIGCOMM 2020)



# Memory Fragmentation

---



# Memory Fragmentation

---



# Memory Fragmentation

---



# Memory Fragmentation

---




# Memory Fragmentation

---



# Memory Fragmentation

---



We want to reorganize or “compact” the allocated regions of memory to be contiguous.



# Memory Fragmentation

---




We want to reorganize or “compact” the allocated regions of memory to be contiguous.

# Virtual Memory – A Quick Primer

---


Modern operating systems maintain a mapping between virtual and physical memory.



# Virtual Memory – A Quick Primer

---


If we relocate objects in physical memory, we have to update their virtual addresses as well.



# Virtual Memory – A Quick Primer

---


But in C and C++, we  
can't alter virtual  
addresses safely.



# Virtual Memory – A Quick Primer

---


How can we relocate objects without changing their virtual addresses?



# Virtual Memory – A Quick Primer

---

We can remap two virtual pages onto the same physical page in memory\*, and discard one of the physical pages.





# Virtual Memory – A Quick Primer

---

We can remap two virtual pages onto the same physical page in memory\*, and discard one of the physical pages.


\* provided there are no collisions between objects on the two pages.






- Memory organized into pages
- Each page holds same # of objects (8)
- Objects are placed on page uniformly at random






We can represent each page as a bitstring, where 0 indicates a free slot and 1 indicates an occupied slot.

We can *mesh* two pages together if they don't have 1s in the same position.




We can represent each page as a bitstring, where 0 indicates a free slot and 1 indicates an occupied slot.


We can *mesh* two pages together if they don't have 1s in the same position.




THEY MESH!



If both bitstrings have a 1 in some position, we can't mesh the strings together.



NO MESH




Now we can forget about the details of memory, and think about our problem in terms of finding meshable pairs of bitstrings.

We want to mesh as many pairs of strings as possible.

1000

0100


1000



0100


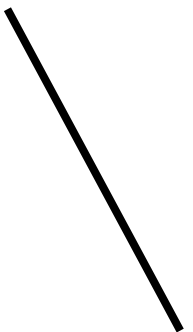
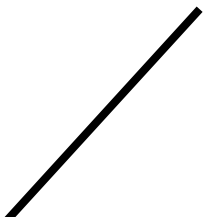


1000




0100

0010




0001


1100



We can think of this as a graph problem!




We can think of this as a graph problem!  
Mesh as many pairs as possible.



# Maximum Matching


---



# Maximum Matching

---

Well-known polynomial time algorithm




# Maximum Matching

---

Well-known polynomial time algorithm

Requires random access to graph



# Maximum Matching


---

Well-known polynomial time algorithm

Requires random access to graph

Do we have random access?

No.





# Not Enough Time!

---

We mesh **during** program execution.

# Not Enough Time!

---

We mesh **during program execution**.

We must “pause” the program to mesh.

# Not Enough Time!

---

We mesh **during program execution**.

We must “pause” the program to mesh.

Only safe to “pause” for a very short time.

# Not Enough Time!


---

We mesh **during program execution**.

We must “pause” the program to mesh.

Only safe to “pause” for a very short time.

Graph doesn't exist yet! We have to do computation to test whether any edge exists, costing valuable time.



# Not Enough Time!

---


We mesh **during program execution**.

We must “pause” the program to mesh.

Only safe to “pause” for a very short time.

Graph doesn't exist yet! We have to do computation to test whether any edge exists, costing valuable time.

Checking an edge is a costly query.



# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

0000100000011000

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

0000100000011000

Say we have a length  $b$  bitstring that's 10% full.

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

0000100000011000

Say we have a length  $b$  bitstring that's 10% full.



# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

$$q = \binom{b - b/10}{b/10} / \binom{b}{b/10}$$

0000100000011000

Say we have a length b bitstring that's 10% full.

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

$$q = \frac{\binom{b - b/10}{b/10}}{\binom{b}{b/10}}$$

In general, if we know how many 1s are in each string, we can get the probability of any two strings meshing.

0000100000011000

Say we have a length b bitstring that's 10% full.

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

$$q = \frac{\binom{b - b/10}{b/10}}{\binom{b}{b/10}}$$

In general, if we know how many 1s are in each string, we can get the probability of any two strings meshing.

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

$$q = \frac{\binom{b - b/10}{b/10}}{\binom{b}{b/10}}$$

In general, if we know how many 1s are in each string, we can get the probability of any two strings meshing.

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

This forms the basis for a greedy query algorithm.

# Random Graphs

---

Recall that the 1s in the bitstrings are distributed randomly.

We can write the probability that it meshes with another 10% full string as

$$q = \frac{\binom{b - b/10}{b/10}}{\binom{b}{b/10}}$$

In general, if we know how many 1s are in each string, we can get the probability of any two strings meshing.

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

This forms the basis for a greedy query algorithm.

But the situation isn't so simple.

# Random Graphs

---

There is a rich body of work on algorithms for random graphs (including matching).

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

This forms the basis for a greedy query algorithm.

But the situation isn't so simple.

# Random Graphs

---

There is a rich body of work on algorithms for random graphs (including matching).

However, this work assumes that the graph's edges are independent.

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

This forms the basis for a greedy query algorithm.

But the situation isn't so simple.

# Random Graphs

---

There is a rich body of work on algorithms for random graphs (including matching).

However, this work assumes that the graph's edges are independent.

That's not the case for meshing graphs!

0000100000011000

We can expect to find a mesh for a bitstring using  $1/q$  queries.

This forms the basis for a greedy query algorithm.

But the situation isn't so simple.



# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

????  
(50% full)


???? (50% full)

???? (50% full)

# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

???? (50% full)  ???? (50% full)

Then if two bitstrings mesh, they must be each other's complement.

???? (50% full)

# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

0110 —— 1001

Then if two bitstrings mesh, they must be each other's complement.

???? (50% full)

# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

0110  1001

Then if two bitstrings mesh, they must be each other's complement.

???? (50% full)

So in this case triangles are impossible. If two of the edges exist, the third must not exist.


# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

Then if two bitstrings mesh, they must be each other's complement.

So in this case triangles are impossible. If two of the edges exist, the third must not exist.




# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

Then if two bitstrings mesh, they must be each other's complement.

So in this case triangles are impossible. If two of the edges exist, the third must not exist.




# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

Then if two bitstrings mesh, they must be each other's complement.

So in this case triangles are impossible. If two of the edges exist, the third must not exist.




# Edge Dependence Example

---

Say we know all of our bitstrings are half full.

Then if two bitstrings mesh, they must be each other's complement.

So in this case triangles are impossible. If two of the edges exist, the third must not exist.

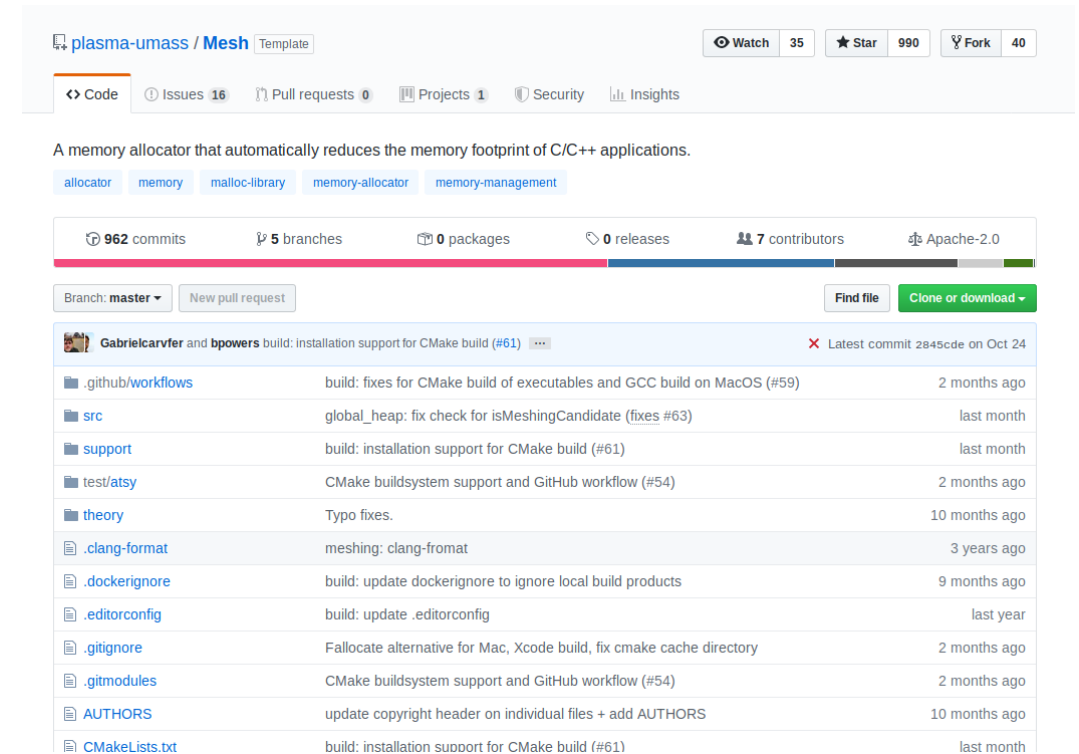


This is a novel and interesting mathematical structure!



# Mesh

We built Mesh, a memory manager powered by a query-limited matching algorithm that can perform memory compaction in C and C++.



plasma-umass / Mesh Template

Watch 35 Star 990 Fork 40

Code Issues 16 Pull requests 0 Projects 1 Security Insights

A memory allocator that automatically reduces the memory footprint of C/C++ applications.

allocator memory malloc-library memory-allocator memory-management

962 commits 5 branches 0 packages 0 releases 7 contributors Apache-2.0

Branch: master New pull request Find file Clone or download

Gabrielcarvfer and bpowers build: installation support for CMake build (#61) Latest commit 2845cde on Oct 24

github/workflows	build: fixes for CMake build of executables and GCC build on MacOS (#59)	2 months ago
src	global_heap: fix check for isMeshingCandidate (fixes #63)	last month
support	build: installation support for CMake build (#61)	last month
test/atsy	CMake buildsystem support and GitHub workflow (#54)	2 months ago
theory	Typo fixes.	10 months ago
.clang-format	meshing: clang-format	3 years ago
.dockerignore	build: update dockerignore to ignore local build products	9 months ago
.editorconfig	build: update .editorconfig	last year
.gitignore	Fallocate alternative for Mac, Xcode build, fix cmake cache directory	2 months ago
.gitmodules	CMake buildsystem support and GitHub workflow (#54)	2 months ago
AUTHORS	update copyright header on individual files + add AUTHORS	10 months ago
CMakeLists.txt	build: installation support for CMake build (#61)	last month

# Mesh

We built Mesh, a memory manager powered by a query-limited matching algorithm that can perform memory compaction in C and C++.

It reduces memory consumption of Firefox by 16% and Redis by 39%.

plasma-umass / Mesh Template

Watch 35 Star 990 Fork 40

Code Issues 16 Pull requests 0 Projects 1 Security Insights

A memory allocator that automatically reduces the memory footprint of C/C++ applications.

allocator memory malloc-library memory-allocator memory-management

962 commits 5 branches 0 packages 0 releases 7 contributors Apache-2.0

Branch: master New pull request Find file Clone or download

Gabrielcarvfer and bpowers build: installation support for CMake build (#61) Latest commit 2845cde on Oct 24

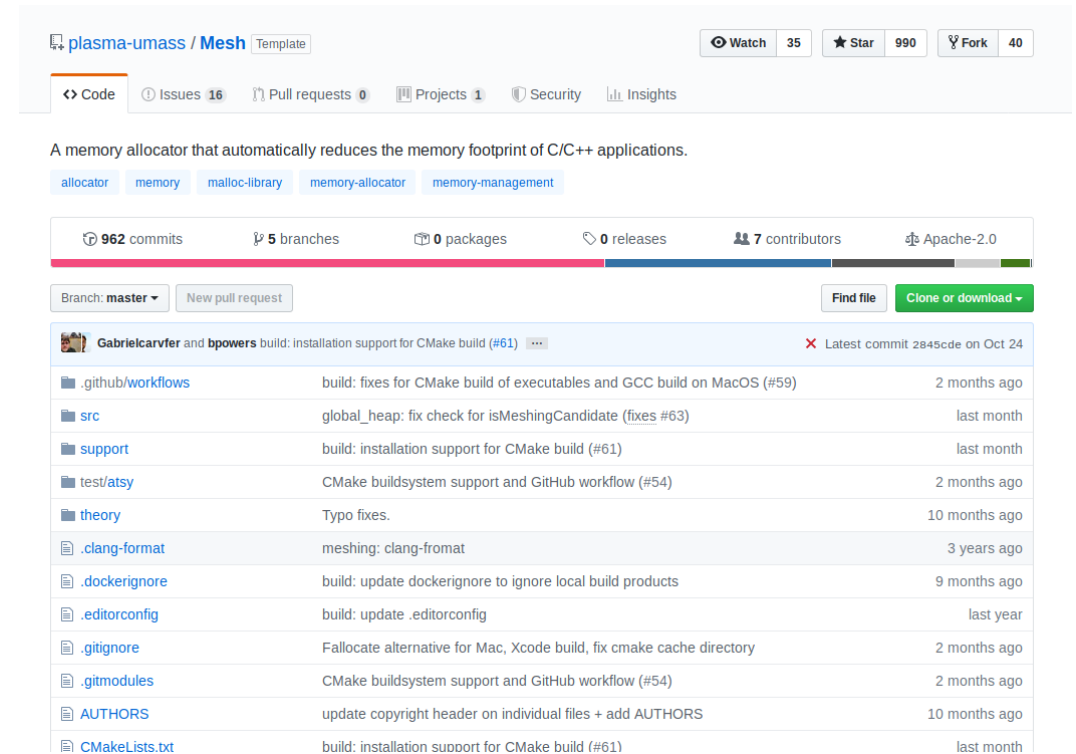
.github/workflows	build: fixes for CMake build of executables and GCC build on MacOS (#59)	2 months ago
src	global_heap: fix check for isMeshingCandidate (fixes #63)	last month
support	build: installation support for CMake build (#61)	last month
test/atsy	CMake buildsystem support and GitHub workflow (#54)	2 months ago
theory	Typo fixes.	10 months ago
.clang-format	meshing: clang-format	3 years ago
.dockerignore	build: update dockerignore to ignore local build products	9 months ago
.editorconfig	build: update .editorconfig	last year
.gitignore	Fallocate alternative for Mac, Xcode build, fix cmake cache directory	2 months ago
.gitmodules	CMake buildsystem support and GitHub workflow (#54)	2 months ago
AUTHORS	update copyright header on individual files + add AUTHORS	10 months ago
CMakeLists.txt	build: installation support for CMake build (#61)	last month

# Mesh

We built Mesh, a memory manager powered by a query-limited matching algorithm that can perform memory compaction in C and C++.

It reduces memory consumption of Firefox by 16% and Redis by 39%.

It is freely available on Github.



plasma-umass / Mesh Template

Watch 35 Star 990 Fork 40

Code Issues 16 Pull requests 0 Projects 1 Security Insights

A memory allocator that automatically reduces the memory footprint of C/C++ applications.

allocator memory malloc-library memory-allocator memory-management

962 commits 5 branches 0 packages 0 releases 7 contributors Apache-2.0


Branch: master New pull request Find file Clone or download

Gabrielcarvfer and bpowers build: installation support for CMake build (#61) Latest commit 2845cde on Oct 24		
github/workflows	build: fixes for CMake build of executables and GCC build on MacOS (#59)	2 months ago
src	global_heap: fix check for isMeshingCandidate (fixes #63)	last month
support	build: installation support for CMake build (#61)	last month
test/atsy	CMake buildsystem support and GitHub workflow (#54)	2 months ago
theory	Typo fixes.	10 months ago
.clang-format	meshing: clang-format	3 years ago
.dockerignore	build: update dockerignore to ignore local build products	9 months ago
.editorconfig	build: update .editorconfig	last year
.gitignore	Fallocate alternative for Mac, Xcode build, fix cmake cache directory	2 months ago
.gitmodules	CMake buildsystem support and GitHub workflow (#54)	2 months ago
AUTHORS	update copyright header on individual files + add AUTHORS	10 months ago
CMakeLists.txt	build: installation support for CMake build (#61)	last month

# Finding the Shape of the Internet


---

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?



# Finding the Shape of the Internet

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?




# Finding the Shape of the Internet

---

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to [colgate.edu](http://colgate.edu)?




# Finding the Shape of the Internet

---

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to [colgate.edu](http://colgate.edu)?

Useful in network measurement.



# Finding the Shape of the Internet


---

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to [colgate.edu](http://colgate.edu)?

Useful in network measurement.

Ex: Can help identify censorship.






# Finding the Shape of the Internet

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to [colgate.edu](http://colgate.edu)?

Useful in network measurement.

Ex: Can help identify censorship.




# Finding the Shape of the Internet

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to [colgate.edu](http://colgate.edu)?

Useful in network measurement.

Ex: Can help identify censorship.




# Finding the Shape of the Internet

I want to visit [www.colgate.edu](http://www.colgate.edu).  
What path through the internet will  
my HTTP request take?

How can I predict paths from *any*  
starting location to colgate.edu?

Measuring a path is a costly query  
called a *traceroute*. We want to  
discover as much of the Internet as  
possible using minimal traceroutes.




# PathCache

A system for efficiently using limited VP measurements to predict paths towards Internet destinations.

Current version discovers 4 times more connections than pre-existing measurement strategies with comparable traceroute budget.

Predicts correct or nearly-correct paths 75% of the time.



---

# Future Work

---

GRAPHS THAT ARE EVEN MORE INCONVENIENT



# Time-Dependent Graph Streams

---




In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading

# Time-Dependent Graph Streams

---




In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading

# Time-Dependent Graph Streams

---




In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading



# Time-Dependent Graph Streams




In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading

# Time-Dependent Graph Streams

---



In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading

# Time-Dependent Graph Streams

---




In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**

What if the order mattered?

Ex: disease spreading

# Time-Dependent Graph Streams

---



In the typical streaming model, graph is the same **regardless of the order edges appear in the stream.**


What if the order mattered?

Ex: disease spreading

Temporal graphs are just beginning to be investigated. No streaming work.

# Time-Dependent Graph Streams

---




Imagine we receive a massive stream of handshakes between many people. Later, we learn one of those people was sick.

Can we determine who is infected without storing the entire stream?

# Time-Dependent Graph Streams


---

This problem is about connectivity or reachability on temporal graphs.



# Time-Dependent Graph Streams

---




This problem is about connectivity or reachability on temporal graphs.

Other potential problems:

# Time-Dependent Graph Streams

---



This problem is about connectivity or reachability on temporal graphs.


Other potential problems:

Determine how long it takes for a disease (or information, or goods) to reach every part of a network.



# Time-Dependent Graph Streams

---



This problem is about connectivity or reachability on temporal graphs.

Other potential problems:


Determine how long it takes for a disease (or information, or goods) to reach every part of a network.

Estimate how many different Patient Zeros could infect a particular person.

# Reconstructing Graphs

---


We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?



# Reconstructing Graphs

---



We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?



# Reconstructing Graphs

---


We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?



# Reconstructing Graphs

---



We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?



# Reconstructing Graphs

---

We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?




# Reconstructing Graphs

---

We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?

More generally, we must reconstruct a matrix  $M$  and can make queries which return a submatrix where rows and columns are deleted randomly. How many queries are needed?




# Reconstructing Graphs

---

We want to reconstruct a graph  $G$ .  
We can make a query which returns a random, unlabeled induced subgraph of  $G$ . How many queries are needed?

More generally, we must reconstruct a matrix  $M$  and can make queries which return a submatrix where rows and columns are deleted randomly. How many queries are needed?

Reconstruction from other queries?





---

Thanks for  
listening!

---